Semantified CEUR-WS in Wikidata

Wolfgang Fahl[1][0000−0002−0821−6995], Tim Holzheim[1,2][0000−0003−2533−6363], Christoph Lange[1,2][0000−0001−9879−3827], Jerven Bolleman[3][0000−0002−7449−1266], and Stefan Decker[1,2][0000−0001−6324−7164]

[1] RWTH Aachen University, Computer Science i5, Aachen, Germany
[2] Fraunhofer FIT, Sankt Augustin, Germany
[3] SIB Swiss Institute of Bioinformatics, Switzerland

Abstract. CEUR Workshop Proceedings (CEUR-WS) has been a non-profit publisher for scientific workshop and conference proceedings since 1995. For more than 30 years it has used HTML and PDF as the single source of truth for the metadata required by indexers such as DBLP and the German national library DNB. Scholars have long craved persistent identifiers for the core entities involved: event series, events, proceedings, papers, editors, authors, institutions, and publishers, with the goal of allowing semantic handling and querying with modern tools.
In this work we show how we enabled the semantification using Wikidata as the target knowledge graph infrastructure. The 2015 Semantic Publishing Challenge queries are finally permanently available with up-to-date metadata. The Scholia portal provides aspect-oriented and user-friendly access and is widely used.
We further point out the challenges of aligning the digital, social, and physical aspects of the ongoing transformation and outline a roadmap to use the available queries to bootstrap a new RDF-based metadata-first publishing approach that will mine metadata from proceedings matter and generate the HTML from it instead of the other way around.

Keywords: Knowledge Graph · Linked Data · Metadata Extraction · Publishing · Semantic Web · Wikidata

1 Introduction

Fig. 1. Most relevant entities for scientific publishing

The Traditional publishing workflows often handle the metadata of the needed core entities (as shown in Figure 1) in an unFAIR way. Findable, Accessible, Interoperable, and Reusable (FAIR) 1.2, Single Point of Truth (SPoT), and Single Source of Truth (SSoT) principles are having low priority. Therefore the opportunities to collect the metadata during the different lifcecyle stages of an

event as shown in table 1 are missed. As technical editorr of CEUR-WS we have been feeling this pain for over a decade. This work shows how the HTML/PDF SSoT of CEUR-WS at the ceur-ws.org domain (SPoT) has been used to create a Knowledge Graph representation of the relevant entities in Wikidata between 2022 and 2026 as a step in this direction thus allowing portals such as Scholia to present the RDF query results in an aspect-oriented and user-friendly way.

Table 1. Scientific Event Lifecycle stages

Stage	**Stakeholders**	**Artifacts Created**	**Format**	**Responsibility**
Announcement	Organizers, Marketing Team	Promotional materials, websites, notifications	HTML, Email	Organizers
Call for Papers (CfP)	Organizers, Authors, Reviewers	CfP documents, emails, submission portals, responses	PDF, HTML, Email	Organizers
Performing Conference	Organizers, Chairs, Presenters, Attendees	Schedules, presentations, participant lists, recordings	PDF, PPT, HTML, Video	Organizers
Proceedings Publication	Editors, Chairs, Authors	Finalized papers, indexed and archived content	PDF, HTML	Editors, Authors, Chairs
Indexing	Libraries, Indexing Services	Catalog entries, metadata records	MARC, XML, RDF	Publishers, Libraries
Referencing/Quoting	Academic Community	Citations, references in future works	BibTeX, HTML	Scholars

Digital Traces of Scientific Events during their lifecyle: Table 1 shows the stages of the scientific event lifecycle that cumulatively create artifacts related to the event.

Moving the responsibility for curating data for relevant entities (see Figure 1) to an earlier stage and thus to the event organizers, while semi-automating the process has several benefits. The quality of the data will increase, and its availability will be much earlier, thus influencing FAIRness—e.g., ensuring computer-readable versions of the core entities are available for downstream systems at each step in the lifecycle. Involving a community such as Wikidata in open source style will further enhance the result [4].

Scientific communication increasingly relies on the use of knowledge graphs, as exemplified by Google Scholar[5] and the evolution of Microsoft Academic Graph [11] into OpenAlex [10]. Legacy systems based on technologies such as MARC [7], PICA [5] or XML are also adapting, highlighted by DBLP's adoption

[4] And will create follow-up issues for the single point of truth data to be discussed later

[5] https://scholar.google.com/

of RDF/SPARQL, with the QLever DBLP SPARQL endpoint[6] [2] being the most up-to-date.

1.1 CEUR Workshop Proceedings

Since 1995 CEUR Workshop Proceedings (CEUR-WS) (https://ceur-ws.org/) has published more than 4,200 Volumes and over 90.000 papers as a free online service for publishing proceedings of scientific workshops (and smaller conferences) in computer science, operated by at team of unpaid volunteers.

Technically, all data and metadata of the proceedings are directly represented as a static website using HTML, PDF and HTTP and FTP protocols as provided by the workshop editors along the CEUR-WS "how to submit" guide [3].

The metadata is available after some delay of weeks or month through indexing services such as dblp [9] and K10plus[12][7] and at the archive of the TIB `https://www.tib.eu/`.

After multiple unfinished attempts to make the metadata of the CEUR-WS platform available for semantic analysis and querying we report on the successful use of wikidata and scholia as the knowledge graph platform to provide FAIR metadata.

Doing so was harder than one would expect - the devil is in the details.

1.2 Making CEUR-WS FAIR

The FAIR principles [8,13] with Persistent Identifiers (PIDs) and rich metadata as core components are being adopted by an increasing number of publishing organizations.

To make CEUR-WS metadata fair a set of relevant queries should be supported, as derived from the original set of queries of the 2014 Semantic publishing challenge as outlined in Section **??**. The metadata should be provided following an ontology that is based on established ontologies. The manual and automatic curation of entries should be possible with public access for all stakeholders, e.g., editors, authors, organizers, publishers and indexers. The infrastructure should be stable and there should be sufficient trust in its long term availability An open source non-profit infrastructure is preferred since this is also the mode of operation of CEUR-WS

Given the HTML/PDF/text input of CEUR-WS, the corresponding Knowledge Graph needs to be created and the single-point-of-truth computer readable metadata be separated from the different representations such as HTML so that the above requirements are fulfilled.

Both the HTML and PDF encoding of the original scientific content are structured for the purpose of optimizing the display / output on paper or screens; therefore, there is a structure loss compared to what was originally available in the text document processor files that the authors might have been using [6].

[6] https://qlever.cs.uni-freiburg.de/dblp
[7] `https://dblp.org/`, `https://opac.k10plus.de`

Most scholars are not aware of this loss in the daily use of published content just because the documents are optimized for display and human consumption [14]. For the metadata extraction and use in knowledge graphs the difference is sometimes disastrous, e.g., when a simple text from a PDF document can not be extracted any more due to exotic styling and formatting or simply because only a scanned graphic image version of an older document is available that needs optical character recognition to extract the textual content see table **??**.

The metadata needed for creating the knowledge graph is only available in natural language/text form and follows rules that have been changed multiple times during the history of CEUR-WS. From 2013 to 2026, there have been 33 different versions of the index file template with no proper tracing of what was changed from version to version. The index and volume HTML files were often edited manually, leading to minor, undocumented differences even within these versions. The usage of the different page versions follows a long-tail Zipfian distribution with 5 versions covering 60% of all volumes.

1.3 Contributions of this Paper

The core contributions presented in this work are:

1. To provide the tools, infrastructure and approaches to modify the CEUR-WS publication workflow to consistently supply high quality FAIR metadata (Semantification) as detailed in section **??**.
2. Splitting the metadata for the three main entities Event, Event Series and Proceedings as a major step towards improving the metadata quality and FAIRness.
3. Providing a bootstrapping [1,4] approach to shift the single point of truth from the current HTML/PDF files to a RDF ready ontology based format.

These contributions are valuable for other publishing use cases see section **??**.

2 From Semantic Publishing Challenge 2015 via Text2KG2023 to Production 2026

3 CEUR-WS Semantification in Wikidata

3.1 Overview

Figure 2 gives an overview of the CEUR-WS Semantification workflow.

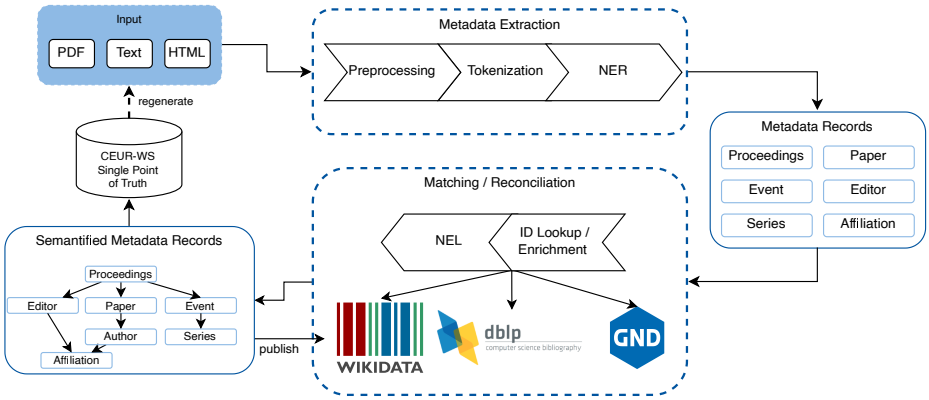

Fig. 2. Workflow of the CEUR-WS Semantification

4 Community Impact

5 Single Point of Truth Bootstrapping

6 Lessons Learned

6.1 Queries as first-class citizens

7 Roadmap: Metadata first Publishing

8 Conclusion

8.1 Supplemental Material Statement

Source code for the prototypes and semantification tools is available via the repositories of the github ceurws organization[8].

8.2 Acknowledgements

We would like to thank Jakob Voß for helping with the k10plus matching and creating the Wikidata property colocated with (P11633)[9] in due time.

Thomas Hoeren and Jonas Kuiter of ITM, Münster kindly gave hints for legal aspects as summarized in section **??**.

This paper is dedicated to the memory of CEUR-WS board member Ralf Klamma † January 2023.

This research has been partly funded by a grant of the Deutsche Forschungsgemeinschaft (DFG). [10]

[8] https://github.com/ceurws
[9] https://www.wikidata.org/wiki/Property:P11633
[10] ConfIDent project; see https://gepris.dfg.de/gepris/projekt/426477583

8.3 Declaration of use of Generative AI

References

1. Bardini, T.: Bootstrapping: Douglas Engelbart, Coevolution, and the Origins of Personal Computing. Stanford University Press, USA (2001)
2. Bast, H., Buchhold, B.: QLever. In: Proceedings of the 2017 ACM Conference on Information and Knowledge Management. pp. 647–656. ACM (Nov 2017). https://doi.org/10.1145/3132847.3132921, https://doi.org/10.1145/3132847.3132921
3. CEUR-WS Editor Team: How to submit to ceur workshop proceedings. CEUR Workshop Proceedings (2019), https://ceur-ws.org/HOWTOSUBMIT.html, online; accessed 2023-12-28
4. Christina Engelbart: About Bootstrapping . https://www.dougengelbart.org/content/view/226/269/ (2007), online; accessed 12 March 2023
5. Costers, L.: The pica catalogue system – paper 26. In: Proceedings of the IATUL Conferences 1979. pp. 73–77. Purdue University (1979), https://docs.lib.purdue.edu/iatul/1979/papers/26/
6. Fahl, W.: The history of scientific publishing (2023), https://cr.bitplan.com/index.php/The_History_of_Scientific_Publishing, online; accessed 2023-03-15
7. Ganseman, J.: Refactoring a library's legacy catalog: a case study. In: IAML Congress 2015. IAML (2015), http://wiki.muziekcollecties.be/images/IAML2015_JG.pdf
8. GO FAIR International Support and Coordination Office: Fair principles (2019), https://www.go-fair.org/fair-principles/
9. Ley, M.: DBLP. Proceedings of the VLDB Endowment **2**(2), 1493–1500 (Aug 2009). https://doi.org/10.14778/1687553.1687577, https://doi.org/10.14778/1687553.1687577
10. Priem, J., Piwowar, H., Orr, R.: OpenAlex: A fully-open index of scholarly works, authors, venues, institutions, and concepts. 26th International Conference on Science, Technology and Innovation Indicators (STI 2022) (2022). https://doi.org/10.5281/ZENODO.6936227, https://zenodo.org/record/6936227
11. Wang, K., Shen, Z., Huang, C., Wu, C.H., Dong, Y., Kanakia, A.: Microsoft Academic Graph: When experts are not enough. Quantitative Science Studies **1**(1), 396–413 (Feb 2020). https://doi.org/10.1162/qss_a_00021, https://doi.org/10.1162/qss_a_00021
12. Wiermann, B.: K10plus – Zehn Bundesländer in einem Bibliothekssystem (2019), https://blog.slub-dresden.de/beitrag/2019/03/27/k10plus-zehn-bundeslaender-in-einem-bibliothekssystem
13. Wilkinson, M.D., Dumontier, M., Aalbersberg, I.J., Appleton, G., Axton, M., Baak, A., Blomberg, N., Boiten, J.W., da Silva Santos, L.B., Bourne, P.E., Bouwman, J., Brookes, A.J., Clark, T., Crosas, M., Dillo, I., Dumon, O., Edmunds, S., Evelo, C.T., Finkers, R., Gonzalez-Beltran, A., Gray, A.J., Groth, P., Goble, C., Grethe, J.S., Heringa, J., 't Hoen, P.A., Hooft, R., Kuhn, T., Kok, R., Kok, J., Lusher, S.J., Martone, M.E., Mons, A., Packer, A.L., Persson, B., Rocca-Serra, P., Roos, M., van Schaik, R., Sansone, S.A., Schultes, E., Sengstag, T., Slater, T., Strawn, G., Swertz, M.A., Thompson, M., van der Lei, J., van Mulligen, E., Velterop, J., Waagmeester, A., Wittenburg, P., Wolstencroft, K., Zhao, J.,

Mons, B.: The FAIR guiding principles for scientific data management and stewardship. Scientific Data **3**(1) (mar 2016). https://doi.org/10.1038/sdata.2016.18, https://doi.org/10.1038/data.2016.18

14. Yu, C., Zhang, C., Wang, J.: Extracting body text from academic PDF documents for text mining. CoRR **abs/2010.12647** (2020), https://arxiv.org/abs/2010.12647